

Launch Decision Memo

FinAssist — Conversational Banking Assistant (Act 1)

Portfolio Case Study Go / No-Go AI Governance

From: Tejaswini Dondeti, AI Safety Reviewer • **To:** AI Governance Review Board (illustrative)
Re: Pre-launch safety decision for FinAssist v2.0 • **Date:** June 2026
Classification: Confidential — fictional institution ("Meridian First Bank")

About this document. The launch decision memo is the artifact a real AI launch-review owner produces — the go/no-go recommendation that converts test evidence into a defensible business decision. Institution fictional; evidence real.

1. Decision

Recommendation: CONDITIONAL GO.

FinAssist v2.0 (hardened) is approved for a **limited, monitored production launch**, subject to the conditions in §4. The one confirmed pre-launch vulnerability has been remediated and re-test is clean; remaining risk is acceptable for a controlled rollout but **not** for an unconditional, unmonitored release.

2. Basis for the decision

Evidence	Finding
Red-team, 150 adversarial tests	Baseline 149/150; hardened 150/150
Overall attack-success rate	0.67% → 0.00% after hardening
Confirmed breach (PI-05, prompt extraction)	Remediated , re-test passes, 0 regressions
Evaluation reliability	Grader required human QC; corrected results sound

The system demonstrates a strong safety posture across all six tested risk classes. The decision is "conditional" rather than "full go" because (a) the suite is static and cannot speak to adaptive attackers, (b) prompt injection remains an industry-wide open problem warranting defense-in-depth, and (c) the evaluation pipeline itself surfaced a reliability finding that must be addressed before scale.

3. Residual risk register (summary)

Risk	Likelihood	Impact	Residual rating
Novel prompt-injection bypass (untested variant)	Medium	High	Medium

Risk	Likelihood	Impact	Residual rating
Confident hallucination of a financial fact	Low	Medium	Low–Medium
PII disclosure via an unforeseen phrasing	Low	High	Medium
Evaluation under-detects a future regression	Medium	High	Medium

Full register with NIST AI 600-1 mappings is provided in the GenAI Risk Assessment.

4. Conditions of approval

- Defense-in-depth.** Deploy an output/content filter for disclosure and PII patterns in addition to the hardened prompt (do not rely on prompt hardening alone).
- Monitored, staged rollout.** Begin with a limited cohort; monitor refusal rates, anomalies, and user-reported issues before broadening.
- Continuous red-teaming.** Stand up a recurring adversarial test cycle; treat the 150-case suite as a regression baseline, not a finish line.
- Grader hardening.** Re-grade with a stronger judge model and add grader-validation sampling before relying on automated evaluation at scale.
- Human escalation + kill switch.** Confirm a working escalation path and a tested rollback/kill-switch procedure prior to launch.
- Scope lock.** Ship the conversational capability only. Any tool-use / money-movement capability is a **separate** launch requiring the Act 2 review (which reaches a stricter verdict).

5. Conditions that would change this decision

This recommendation would downgrade to **NO-GO** if, before launch: continuous red-teaming surfaces an unremediated High-severity breach; the defense-in-depth filter cannot be deployed; or the production model differs materially from the evaluated Llama 3.3 70B without re-validation.

6. Sign-off

Role	Name	Decision	Date
AI Safety Reviewer (author)	Tejaswini Dondeti	Conditional Go	June 2026
Governance Review Board	<i>(illustrative)</i>	<i>pending</i>	—
Product Owner	<i>(illustrative)</i>	<i>pending</i>	—

This memo should be read together with the Red-Team Evaluation Report, System Card, GenAI Risk Assessment, and Framework Crosswalk.